



# Breast Cancer Classification with a simple Neural Network (NN)

## Importing the Dependencies

```
In [3]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report, confusion_matrix
import tensorflow as tf
from tensorflow import keras
import matplotlib.pyplot as plt
import seaborn as sns
```

## Data Collection & Processing

```
In [4]: df = pd.read_csv("data.csv")
df.drop("id", axis=1, inplace=True)
df['diagnosis'] = df['diagnosis'].map({'M':0, 'B':1})
```

```
In [5]: df
```

```
Out[5]:
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoo
0	0	17.99	10.38	122.80	1001.0	
1	0	20.57	17.77	132.90	1326.0	
2	0	19.69	21.25	130.00	1203.0	
3	0	11.42	20.38	77.58	386.1	
4	0	20.29	14.34	135.10	1297.0	
...	...	...	...	...	...	
564	0	21.56	22.39	142.00	1479.0	
565	0	20.13	28.25	131.20	1261.0	
566	0	16.60	28.08	108.30	858.1	
567	0	20.60	29.33	140.10	1265.0	
568	1	7.76	24.54	47.92	181.0	

569 rows x 31 columns

```
In [6]: # print last 5 rows of the dataframe
df.tail()
```

```
Out[6]:
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoo
<b>564</b>	0	21.56	22.39	142.00	1479.0	
<b>565</b>	0	20.13	28.25	131.20	1261.0	
<b>566</b>	0	16.60	28.08	108.30	858.1	
<b>567</b>	0	20.60	29.33	140.10	1265.0	
<b>568</b>	1	7.76	24.54	47.92	181.0	

5 rows × 31 columns

```
In [7]: # number of rows and columns in the dataset
df.shape
```

```
Out[7]: (569, 31)
```

```
In [8]: # getting some information about the data
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   diagnosis                             569 non-null    int64
1   radius_mean                           569 non-null    float64
2   texture_mean                           569 non-null    float64
3   perimeter_mean                         569 non-null    float64
4   area_mean                             569 non-null    float64
5   smoothness_mean                        569 non-null    float64
6   compactness_mean                       569 non-null    float64
7   concavity_mean                         569 non-null    float64
8   concave points_mean                    569 non-null    float64
9   symmetry_mean                          569 non-null    float64
10  fractal_dimension_mean                 569 non-null    float64
11  radius_se                              569 non-null    float64
12  texture_se                              569 non-null    float64
13  perimeter_se                           569 non-null    float64
14  area_se                                569 non-null    float64
15  smoothness_se                          569 non-null    float64
16  compactness_se                         569 non-null    float64
17  concavity_se                           569 non-null    float64
18  concave points_se                      569 non-null    float64
19  symmetry_se                            569 non-null    float64
20  fractal_dimension_se                   569 non-null    float64
21  radius_worst                           569 non-null    float64
22  texture_worst                           569 non-null    float64
23  perimeter_worst                        569 non-null    float64
24  area_worst                             569 non-null    float64
25  smoothness_worst                       569 non-null    float64
26  compactness_worst                      569 non-null    float64
27  concavity_worst                        569 non-null    float64
28  concave points_worst                   569 non-null    float64
29  symmetry_worst                         569 non-null    float64
30  fractal_dimension_worst                 569 non-null    float64
dtypes: float64(30), int64(1)
memory usage: 137.9 KB

```

```

In [9]: # checking for missing values
df.isnull().sum()

```

Out[9]:

0

<b>diagnosis</b>	0
<b>radius_mean</b>	0
<b>texture_mean</b>	0
<b>perimeter_mean</b>	0
<b>area_mean</b>	0
<b>smoothness_mean</b>	0
<b>compactness_mean</b>	0
<b>concavity_mean</b>	0
<b>concave points_mean</b>	0
<b>symmetry_mean</b>	0
<b>fractal_dimension_mean</b>	0
<b>radius_se</b>	0
<b>texture_se</b>	0
<b>perimeter_se</b>	0
<b>area_se</b>	0
<b>smoothness_se</b>	0
<b>compactness_se</b>	0
<b>concavity_se</b>	0
<b>concave points_se</b>	0
<b>symmetry_se</b>	0
<b>fractal_dimension_se</b>	0
<b>radius_worst</b>	0
<b>texture_worst</b>	0
<b>perimeter_worst</b>	0
<b>area_worst</b>	0
<b>smoothness_worst</b>	0
<b>compactness_worst</b>	0
<b>concavity_worst</b>	0
<b>concave points_worst</b>	0
<b>symmetry_worst</b>	0
<b>fractal_dimension_worst</b>	0

**dtype:** int64


```
In [10]: # statistical measures about the data
df.describe()
```

```
Out[10]:
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	si
<b>count</b>	569.000000	569.000000	569.000000	569.000000	569.000000	
<b>mean</b>	0.627417	14.127292	19.289649	91.969033	654.889104	
<b>std</b>	0.483918	3.524049	4.301036	24.298981	351.914129	
<b>min</b>	0.000000	6.981000	9.710000	43.790000	143.500000	
<b>25%</b>	0.000000	11.700000	16.170000	75.170000	420.300000	
<b>50%</b>	1.000000	13.370000	18.840000	86.240000	551.100000	
<b>75%</b>	1.000000	15.780000	21.800000	104.100000	782.700000	
<b>max</b>	1.000000	28.110000	39.280000	188.500000	2501.000000	

8 rows × 31 columns

```
In [11]: # Class distribution
sns.countplot(x='diagnosis', data=df)
plt.title("Target Class Distribution")
plt.show()
```



```
In [12]: # checking the distribution of Target Variable  
df['diagnosis'].value_counts()
```

```
Out[12]:
```

	count
--	-------

diagnosis	
1	357
0	212

**dtype:** int64

1 --> Benign

0 --> Malignant

```
In [13]: df.groupby('diagnosis').mean()
```

Out[13]:

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothnes
diagnosis					
0	17.462830	21.604906	115.365377	978.376415	0
1	12.146524	17.914762	78.075406	462.790196	0

2 rows × 30 columns

Separating the features and target

```
In [14]: X = df.drop(columns='diagnosis', axis=1)
Y = df['diagnosis']
```

```
In [15]: print(X)
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	\
0	17.99	10.38	122.80	1001.0	0.11840	
1	20.57	17.77	132.90	1326.0	0.08474	
2	19.69	21.25	130.00	1203.0	0.10960	
3	11.42	20.38	77.58	386.1	0.14250	
4	20.29	14.34	135.10	1297.0	0.10030	
..	...	...	...	...	...	
564	21.56	22.39	142.00	1479.0	0.11100	
565	20.13	28.25	131.20	1261.0	0.09780	
566	16.60	28.08	108.30	858.1	0.08455	
567	20.60	29.33	140.10	1265.0	0.11780	
568	7.76	24.54	47.92	181.0	0.05263	

	compactness_mean	concavity_mean	concave points_mean	symmetry_mean	\
0	0.27760	0.30010	0.14710	0.2419	
1	0.07864	0.08690	0.07017	0.1812	
2	0.15990	0.19740	0.12790	0.2069	
3	0.28390	0.24140	0.10520	0.2597	
4	0.13280	0.19800	0.10430	0.1809	
..	...	...	...	...	
564	0.11590	0.24390	0.13890	0.1726	
565	0.10340	0.14400	0.09791	0.1752	
566	0.10230	0.09251	0.05302	0.1590	
567	0.27700	0.35140	0.15200	0.2397	
568	0.04362	0.00000	0.00000	0.1587	

	fractal_dimension_mean	...	radius_worst	texture_worst	\
0	0.07871	...	25.380	17.33	
1	0.05667	...	24.990	23.41	
2	0.05999	...	23.570	25.53	
3	0.09744	...	14.910	26.50	
4	0.05883	...	22.540	16.67	
..	...	...	...	...	
564	0.05623	...	25.450	26.40	
565	0.05533	...	23.690	38.25	
566	0.05648	...	18.980	34.12	
567	0.07016	...	25.740	39.42	
568	0.05884	...	9.456	30.37	

	perimeter_worst	area_worst	smoothness_worst	compactness_worst	\
0	184.60	2019.0	0.16220	0.66560	
1	158.80	1956.0	0.12380	0.18660	
2	152.50	1709.0	0.14440	0.42450	
3	98.87	567.7	0.20980	0.86630	
4	152.20	1575.0	0.13740	0.20500	
..	...	...	...	...	
564	166.10	2027.0	0.14100	0.21130	
565	155.00	1731.0	0.11660	0.19220	
566	126.70	1124.0	0.11390	0.30940	
567	184.60	1821.0	0.16500	0.86810	
568	59.16	268.6	0.08996	0.06444	

	concavity_worst	concave points_worst	symmetry_worst	\
0	0.7119	0.2654	0.4601	



1	0.2416	0.1860	0.2750
2	0.4504	0.2430	0.3613
3	0.6869	0.2575	0.6638
4	0.4000	0.1625	0.2364
..	...	...	...
564	0.4107	0.2216	0.2060
565	0.3215	0.1628	0.2572
566	0.3403	0.1418	0.2218
567	0.9387	0.2650	0.4087
568	0.0000	0.0000	0.2871

	fractal_dimension_worst
0	0.11890
1	0.08902
2	0.08758
3	0.17300
4	0.07678
..	...
564	0.07115
565	0.06637
566	0.07820
567	0.12400
568	0.07039

[569 rows x 30 columns]

In [16]: `print(Y)`

```
0      0
1      0
2      0
3      0
4      0
..
564    0
565    0
566    0
567    0
568    1
Name: diagnosis, Length: 569, dtype: int64
```

Splitting the data into training data & Testing data

In [17]: `# Splitting the dataset into training and testing sets`  
`X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=42)`

In [18]: `print(X.shape, X_train.shape, X_test.shape)`


(569, 30) (455, 30) (114, 30)

Standardize the data

In [19]: `scaler = StandardScaler()`

```
X_train_std = scaler.fit_transform(X_train)
X_test_std = scaler.transform(X_test)
```

## Building the Neural Network



```
In [20]: tf.random.set_seed(3)
```

```
In [21]: # setting up the layers of Neural Network
```

```
# Defining the Neural Network model
```

```
model = keras.Sequential([
    keras.layers.Dense(20, activation='relu', input_shape=(30,)),
    keras.layers.Dense(1, activation='sigmoid')
])
```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
In [22]: # compiling the Neural Network
```

```
model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=['accuracy'])
```

```
In [23]: # training the Neural Network
```

```
# Training the model
```

```
history = model.fit(X_train_std, Y_train, validation_split=0.1, epochs=10)
```

Epoch 1/10  
13/13 ————— 2s 34ms/step - accuracy: 0.3723 - loss: 1.0333 - val\_accuracy: 0.4565 - val\_loss: 0.7439  
Epoch 2/10  
13/13 ————— 0s 14ms/step - accuracy: 0.4782 - loss: 0.7535 - val\_accuracy: 0.8043 - val\_loss: 0.5189  
Epoch 3/10  
13/13 ————— 0s 13ms/step - accuracy: 0.6648 - loss: 0.5571 - val\_accuracy: 0.9783 - val\_loss: 0.3784  
Epoch 4/10  
13/13 ————— 0s 7ms/step - accuracy: 0.8087 - loss: 0.4306 - val\_accuracy: 0.9783 - val\_loss: 0.2954  
Epoch 5/10  
13/13 ————— 0s 7ms/step - accuracy: 0.8535 - loss: 0.3507 - val\_accuracy: 0.9783 - val\_loss: 0.2445  
Epoch 6/10  
13/13 ————— 0s 7ms/step - accuracy: 0.8787 - loss: 0.2975 - val\_accuracy: 0.9783 - val\_loss: 0.2115  
Epoch 7/10  
13/13 ————— 0s 7ms/step - accuracy: 0.9021 - loss: 0.2597 - val\_accuracy: 0.9783 - val\_loss: 0.1885  
Epoch 8/10  
13/13 ————— 0s 7ms/step - accuracy: 0.9234 - loss: 0.2311 - val\_accuracy: 0.9783 - val\_loss: 0.1715  
Epoch 9/10  
13/13 ————— 0s 11ms/step - accuracy: 0.9282 - loss: 0.2086 - val\_accuracy: 0.9783 - val\_loss: 0.1585  
Epoch 10/10  
13/13 ————— 0s 7ms/step - accuracy: 0.9405 - loss: 0.1903 - val\_accuracy: 0.9783 - val\_loss: 0.1482


Visualizing accuracy and loss

```
In [24]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])

plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')

plt.legend(['training data', 'validation data'], loc = 'lower right')
```

Out[24]: <matplotlib.legend.Legend at 0x7ab114d867d0>




```
In [25]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])

plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')

plt.legend(['training data', 'validation data'], loc = 'upper right')
```

Out[25]: <matplotlib.legend.Legend at 0x7ab176feba50>



Accuracy of the model on test data

```
In [26]: # Evaluating the model on test data
loss, accuracy = model.evaluate(X_test_std, Y_test)
print(accuracy)
```

4/4 ————— 0s 11ms/step - accuracy: 0.9230 - loss: 0.2199  
0.9298245906829834

```
In [27]: print(X_test_std.shape)
print(X_test_std[0])
```

```
(114, 30)
[-0.04462793 -1.41612656 -0.05903514 -0.16234067  2.0202457  -0.11323672
  0.18500609  0.47102419  0.63336386  0.26335737  0.53209124  2.62763999
  0.62351167  0.11405261  1.01246781  0.41126289  0.63848593  2.88971815
 -0.41675911  0.74270853 -0.32983699 -1.67435595 -0.36854552 -0.38767294
  0.32655007 -0.74858917 -0.54689089 -0.18278004 -1.23064515 -0.6268286 ]
```

```
In [28]: Y_pred = model.predict(X_test_std)
```

4/4 ————— 0s 16ms/step

```
In [29]: print(Y_pred.shape)
print(Y_pred[0])
```

```
(114, 1)
[0.57944816]
```

```
In [30]: print(X_test_std)
```

```
[[ -0.04462793 -1.41612656 -0.05903514 ... -0.18278004 -1.23064515
   -0.6268286 ]
 [  0.24583601 -0.06219797  0.21802678 ...  0.54129749  0.11047691
   0.0483572 ]
 [-1.26115925 -0.29051645 -1.26499659 ... -1.35138617  0.269338
  -0.28231213]
 ...
 [  0.72709489  0.45836817  0.75277276 ...  1.46701686  1.19909344
   0.65319961]
 [  0.25437907  1.33054477  0.15659489 ... -1.29043534 -2.22561725
  -1.59557344]
 [  0.84100232 -0.06676434  0.8929529  ...  2.15137705  0.35629355
   0.37459546]]
```

```
In [31]: print(Y_pred)
```

[5.7944816e-01]  
[5.4799771e-01]  
[9.5382488e-01]  
[6.0756531e-05]  
[4.8952186e-01]  
[7.6010404e-03]  
[6.2022758e-01]  
[9.7304952e-01]  
[9.0360391e-01]  
[9.3224633e-01]  
[6.2372327e-01]  
[8.6918908e-01]  
[6.6956234e-01]  
[8.4866720e-01]  
[8.8146442e-01]  
[1.0968794e-02]  
[9.5920968e-01]  
[8.5751635e-01]  
[8.6840624e-01]  
[2.2778651e-02]  
[1.6535509e-01]  
[9.5268559e-01]  
[9.2469513e-01]  
[9.6149290e-01]  
[7.5835764e-01]  
[2.8978478e-02]  
[8.8868731e-01]  
[7.7837563e-01]  
[2.1731045e-02]  
[1.4245502e-02]  
[8.7887466e-01]  
[8.6644804e-01]  
[8.6391968e-01]  
[5.1455584e-04]  
[1.6218390e-02]  
[6.7806208e-01]  
[9.8855454e-01]  
[8.9672464e-01]  
[9.6932524e-01]  
[9.0333766e-01]  
[5.4680562e-04]  
[1.6645780e-01]  
[9.8668027e-01]  
[8.0902153e-01]  
[1.5716387e-01]  
[8.9677358e-01]  
[9.3963349e-01]  
[8.7662619e-01]  
[3.8213830e-04]  
[4.3045968e-02]  
[9.1671801e-01]  
[1.4556846e-01]  
[4.5559421e-01]  
[9.4883031e-01]

[9.2862117e-01]  
[4.7570056e-01]  
[9.1469473e-01]  
[8.9343512e-01]  
[1.6958913e-02]  
[7.3393339e-01]  
[7.0428139e-01]  
[7.2126172e-02]  
[9.5788360e-01]  
[2.3276834e-02]  
[2.6675813e-02]  
[4.8119459e-01]  
[1.9572754e-03]  
[6.3562877e-02]  
[5.7100791e-01]  
[9.5968835e-02]  
[1.2510759e-01]  
[1.5738681e-02]  
[8.3469772e-01]  
[1.7510022e-01]  
[9.8128551e-01]  
[1.4944032e-01]  
[8.9824814e-01]  
[9.2390132e-01]  
[5.2017534e-01]  
[2.3615752e-01]  
[1.0397666e-02]  
[1.3247216e-01]  
[3.2603938e-02]  
[5.7312453e-01]  
[8.5071766e-01]  
[4.5505381e-01]  
[8.3423293e-01]  
[7.5447232e-01]  
[5.6370598e-01]  
[1.0403588e-02]  
[8.6998773e-01]  
[9.0814209e-01]  
[8.2195503e-01]  
[1.9545414e-02]  
[1.3483056e-01]  
[8.4950876e-01]  
[7.0937891e-03]  
[2.3816984e-02]  
[9.4713300e-01]  
[9.7018409e-01]  
[9.7874528e-01]  
[3.5689881e-01]  
[2.3128826e-04]  
[1.3958461e-03]  
[9.0535390e-01]  
[9.0754831e-01]  
[8.0991483e-01]  
[7.5048649e-01]



```
[9.3533069e-01]
[9.2685550e-01]
[2.8191048e-03]
[1.0194858e-02]
[5.1812589e-01]
[4.5157094e-02]
```

`model.predict()` gives the prediction probability of each class for that data point

```
In [32]: # argmax function
```

```
my_list = [0.25, 0.56]

index_of_max_value = np.argmax(my_list)
print(my_list)
print(index_of_max_value)
```

[0.25, 0.56]  
1

```
In [33]: # converting the prediction probability to class labels
```

```
Y_pred_labels = [np.argmax(i) for i in Y_pred]
print(Y_pred_labels)
```

[illegible]

## Building the predictive system

```
In [34]: def predict_cancer(input_data, model, scaler):
```

|| || ||

Predicts whether a tumor is benign or malignant based on input features.

### Parameters:

- input\_data: Tuple or list of 30 numerical values (features)
- model: Trained Keras model

- scaler: Trained StandardScaler instance

Returns:

- String indicating the result: 'Malignant' or 'Benign'

"""

*# Convert input to numpy array*

input\_data\_as\_numpy\_array = np.asarray(input\_data)

*# Reshape as 2D array (single sample)*

input\_data\_resaped = input\_data\_as\_numpy\_array.reshape(1, -1)

*# Standardize input*

input\_data\_std = scaler.transform(input\_data\_resaped)

*# Predict*

prediction = model.predict(input\_data\_std)

*# Binary classification threshold*

predicted\_class = int(prediction[0][0] > 0.5)

*# Return result*

return 'Malignant' if predicted\_class == 0 else 'Benign'

```
In [35]: input_data = (11.76,21.6,74.72,427.9,0.08637,0.04966,0.01657,0.01115,0.1495,0.
          0.4062,1.21,2.635,28.47,0.005857,0.009758,0.01168,0.007445,0.024
          12.98,25.72,82.98,516.5,0.1085,0.08615,0.05523,0.03715,0.2433,0.

result = predict_cancer(input_data, model, scaler)
print("Prediction:", result)
```

1/1 ————— 0s 38ms/step

Prediction: Benign

/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names  
warnings.warn(